

LAB6 REPORT - COPY ON WRITE

CS19B043 - T. BHANU SHASHANK

1)

- a) Modified the uvmcopy() to map the parent's physical pages to child
- b) Cleared the PTE_W Bits in both parent and child
- c) Finally setting PTE_C bits.

Note: RSW (1<= 8) and PCOUNTACCESS(va) are defined in riscv.h file.

Code : (vm.c file)

```
int
uvmcopy(pagetable_t old, pagetable_t new, uint64 sz){
    pte_t *pte;
    uint64 pa, i;
    uint flags;
    for(i = 0; i < sz; i += PGSIZE){                //For each page frame.
        if((pte = walk(old, i, 0)) == 0)
            panic("uvmcopy: pte should exist");
        if((*pte & PTE_V) == 0)
            panic("uvmcopy: page not present");
        pa = PTE2PA(*pte);
        flags = PTE_FLAGS(*pte);
        flags &= ~(PTE_W);                          // clear the write bit
        flags |= RSW;                                // set the RSW bits to indicate a COW page
        *pte = PA2PTE(pa) | flags;
        if(mappages(new, i, PGSIZE, pa, flags) != 0)
        {
            printf("page mapping error in old\n");
            goto err;
        }
        pcount[PCOUNTACCESS(pa)]++;                  //increasing pcount to store number of owner's
    }
    return 0;
err:
    uvmunmap(new, 0, i, 1);
    return -1;
}
```

Explanation :

- a) For each page frame of both parent and child, clear the write bit.
- b) Set the RSW bit to indicate it is Copy on the write page.
- c) Increasing the pcount to store the number of owner's.

2) Modified usertrap() to recognize page faults. When a page-fault occurs on a COW page, allocate a new page with kalloc(), copy the old page to the new page, and install the new page in the PTE with PTE_W set in trap.c file.

Code : (usertrap() in trap.c file)

```

else if(r_scause() == 15){                                     // Store/AMO page fault
    uint64 faulting_va, pa;
    pte_t *pte; uint flags; char *mem;
    faulting_va = r_stval();
    if((faulting_va >= MAXVA) || ((pte = walk2(p->pagetable, faulting_va, 0)) == 0)){
        printf("page mapping does not exist\n");
        p->killed = 1;
        exit(-1); }
    if(((pte & PTE_V) == 0) || ((pte & PTE_U) == 0)){
        printf("page mapping does not exist\n");
        p->killed = 1;
        exit(-1); }
    pa = PTE2PA(*pte);
    flags = PTE_FLAGS(*pte);
    if(*pte & RSW) { // COW page check.
        if(pcount[PCOUNTACCESS(pa)] == 1){ //If there is single owner then just set the write bit
            *pte &= ~(RSW);
            *pte |= PTE_W;
        }
        else{
            if((mem = kalloc()) == 0){ //allocate memory
                printf("no free pages\n");
                p->killed = 1;
                exit(-1);
            }
            memmove(mem, (char *)pa, PGSIZE); //copy old page to new page
            flags &= ~(RSW); //clear RSW bits
            flags |= PTE_W; //set write bit
            uvmunmap(p->pagetable, PGROUNDDOWN(faulting_va), 1, 1);
            if(mappages(p->pagetable, PGROUNDDOWN(faulting_va), PGSIZE, (uint64)mem, flags) != 0){
                kfree(mem);
                printf("page mapping error\n");
                p->killed = 1;
                exit(-1); }
        }
    }
    else {
        printf("invalid page fault\n");
        p->killed = 1;
        exit(-1);}
}

```

Explanation:

- a) Which type of interrupt is stored in the scause register if scause value is 15 (page-fault).
- b) Checking if there is a single owner of a COW page, then it's cow and non-writable flags can be reset and it need not obtain a copy of the page.
- c) Else allocate new page and copy old page to new page with set PTE_W bit in PTE.
- d) Followed by sanity checks.

3) Each physical page is free-ed when the last PTE reference to it goes away, by implementing reference counts in kalloc.c.

Code:

```
extern int pcount[PCOUNTACCESS(PHYSTOP)];
void kfree(void *pa){
    struct run *r;
    if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
        panic("kfree");
    if(((uint64)pa >= KERNBASE) && pcount[PCOUNTACCESS((uint64)pa)] > 0) {
        pcount[PCOUNTACCESS((uint64)pa)]--;    // Decreasing the owner's by one
        if(pcount[PCOUNTACCESS((uint64)pa)] > 0) {
            return;
        }
    }
    // Free-up the page if it last PTE reference.
    memset(pa, 1, PGSIZE);                // Fill with junk to catch dangling refs.
    r = (struct run*)pa;
    acquire(&kmem.lock);
    r->next = kmem.freelist;
    kmem.freelist = r;
    release(&kmem.lock);
}

void *kalloc(void){
    struct run *r;
    acquire(&kmem.lock);
    r = kmem.freelist;
    if(r)
        kmem.freelist = r->next;
    release(&kmem.lock);
    if(r){
        if((uint64)r >= KERNBASE) {
            pcount[PCOUNTACCESS((uint64)r)] = 1;    // Intialize the pcount with one
        }
        memset((char*)r, 5, PGSIZE); // fill with junk
    }
    return (void*)r;}

```

Explanation:

- a) Checking if it is the last PTE reference or not.
- b) If it is the last PTE reference then it is free-ed and added to kfree list.
- c) Else don't change anything.
- d) And second function kalloc initializes the pcount to one, since it is allocating first time.

4) Modified copyout() to use the same scheme as page faults when it encounters a COW page.

Code: (additional code in copyout function in vm.c file)

```

if((va0 >= MAXVA) || ((pte = walk(pagetable, va0, 0)) == 0)){
    printf("invalid page\n");
    return -1;
}
if(((pte & PTE_V) == 0) || ((pte & PTE_U) == 0)){
    printf("invalid page\n");
    return -1;
}
pa0 = PTE2PA(*pte);
flags = PTE_FLAGS(*pte);
if(*pte & RSW) {
    if((mem = kalloc()) == 0)
    {
        printf("no free pages\n");
        return -1;
    }
    memmove(mem, (char *)pa0, PGSIZE);
    flags &= ~(RSW); //clear RSW bits
    flags |= PTE_W; //set write bit
    uvmunmap(pagetable, va0, 1, 1);
    if(mappages(pagetable, va0, PGSIZE, (uint64)mem, flags) != 0)
    {
        kfree(mem);
        printf("page mapping error\n");
        return -1;
    }
}

```

// COW page
// allocating memory

// no free pages present.

Explanation:

- a) Checking where it is COW page or not using the RSW bit.
- b) Allocate the new page, copy the old page and set the flag bit such that it is writable (PTE_W).
- c) There are some sanity checks in between.
- d) Here implementation is similar to page-fault handled in user trap() in trap.c.